

Team 17

107060008 張承勛

107062123 鄭宇軒

107062315 黃偉哲

Implementation

First and foremost, add the keyword “explain” into the keyword-list in

`query/parse/Lexer` so that the lexer can recognize this new keyword. Then, check if the keyword “explain” is contained in the query inside `query/parse/Parser`; if so, eat it and pass the information to `QueryData`.

Next, create a new class as `ExplainPlan.java` and call it in the planner if we need to explain the query. By the specification of this homework, JDBC client will get the result through `RemoteResultSet.getString("query-plan")`, therefore we need to add a field called “query-plan” in the schema. Besides, add a method called `explainStr` in every plan so that each plan can explain its own estimations, and then we can call the method in `ExplainPlan` so that it can be operated recursively till the `TablePlan`. Collect all the information and pass it to `ExplainScan`. Thus, we also need to create a new class as `ExplainScan.java` for `ExplainPlan.java` to scan the output for `ExplainPlan` to employ. Again, by the specification of this homework, we should pass the information about those estimations from each plan to the method `getVal(String fldName)` inside `ExplainScan` so that the JDBC client can get the output successfully.

As for the number of actually accessed records, our implementation is to check how many `s.next()` can be called and return `TRUE` in `ExplainPlan`.

Experiment

- The example query

```
SQL> EXPLAIN SELECT COUNT(d_id) FROM district, warehouse WHERE d_w_id = w_id GROUP BY w_id
query-plan
-----
->ProjectPlan (#blks=2, #recs=1)
  ->GroupByPlan: (#blks=2, #recs=1)
    ->SortPlan (#blks=2, #recs=10)
      ->SelectPlan pred:(d_w_id=w_id) (#blks=22, #recs=10)
        ->ProductPlan (#blks=22, #recs=10)
          ->TablePlan on (warehouse) (#blks=2, #recs=1)
          ->TablePlan on (district) (#blks=2, #recs=10)

Actual #recs: 1
```

- A query accessing single table with WHERE

```
SQL> EXPLAIN SELECT i_name FROM item WHERE i_id < 87

query-plan
-----
->ProjectPlan (#blks=6251, #recs=5)
    ->SelectPlan pred:(i_id<87.0) (#blks=6251, #recs=5)
        ->TablePlan on (item) (#blks=6251, #recs=100000)

Actual #recs: 86
```

In this experiment, it is clearly that the actual #recs is correct.

- A query accessing multiple tables with WHERE

```
SQL> EXPLAIN SELECT w_id FROM warehouse, district WHERE w_id = d_w_id AND w_id < 10

query-plan
-----
->ProjectPlan (#blks=22, #recs=10)
    ->SelectPlan pred:(w_id=d_w_id and w_id<10.0) (#blks=22, #recs=10)
        ->ProductPlan (#blks=22, #recs=10)
            ->TablePlan on (warehouse) (#blks=2, #recs=1)
            ->TablePlan on (district) (#blks=2, #recs=10)

Actual #recs: 10
```

- A query with ORDER BY

```
SQL> EXPLAIN SELECT i_id FROM item WHERE i_id < 15 ORDER BY i_id DESC

query-plan
-----
->SortPlan (#blks=0, #recs=0)
    ->ProjectPlan (#blks=6251, #recs=0)
        ->SelectPlan pred:(i_id<15.0) (#blks=6251, #recs=0)
            ->TablePlan on (item) (#blks=6251, #recs=100000)

Actual #recs: 14
```

- A query with GROUP BY and at least one aggregation function

```
SQL> EXPLAIN SELECT c_id, SUM(c_discount) FROM customer WHERE c_id < 10 GROUP BY c_id

query-plan
-----
->ProjectPlan (#blks=36, #recs=2)
    ->GroupByPlan: (#blks=36, #recs=2)
        ->SortPlan (#blks=36, #recs=71)
            ->SelectPlan pred:(c_id<10.0) (#blks=15001, #recs=71)
                ->TablePlan on (customer) (#blks=15001, #recs=30000)

Actual #recs: 9
```